



ALPHA OMEGA

| | Creating New Possibilities

U.S.
SEC

▲
AWS
PAiTH

⚡
Continuum
Design

Zscaler +
SPLX.ai

SECURE ENTERPRISE AI · PROOF-OF-VALUE PROGRAM

Deploying Trusted AI for the SEC on the Private AI Tech Hub

Mission-aligned AI Proof-of-Value demos for Enforcement, DERA, and Examinations — built fast with Continuum Design, deployed on Amazon Bedrock inside the secure AWS PAiTH landing zone, and protected with Zscaler zero-trust and SPLX.ai red-teaming.

PLATFORM

AWS PAiTH · Amazon Bedrock

BUILD ACCELERATOR

Continuum Design

SECURITY

Zscaler + SPLX.ai

What we'll walk through **today**

1 PAiTH

The secure, governed home for federal generative AI.

3 SEC "Fed Claw" Fraud Monitor

Enforcement — asset recovery before funds go dark.

5 Zscaler + SPLX

Zero-trust access plus continuous AI red-teaming.

2 Continuum Design

How we turn ideas into working systems in hours.

4 Agentic Examiner

EXAMS — pre-exam intelligence, ready to demo live.

6 XBRL Quality Guard

DERA — data integrity, proven in 2 weeks.

Closing: the consolidated business value across every initiative.

AWS PAiTH & AI — a secure home for federal AI

The Private AI Tech Hub lets agencies adopt commercial-grade generative AI without surrendering control of their data, models, or compliance posture.

Secure by design

- Isolated VPC landing zone, separate from production systems
- Zero-trust access & label-based access control
- FIPS 140-2 validated secret management
- AWS Cognito authentication at the edge


Enterprise model choice

- Amazon Bedrock — enterprise-approved foundation models
- Swap or compare models without re-architecting apps
- Plugins & agents layer for custom workflows
- Knowledge bases & files ingestor for grounding

Governed & observable

- Bedrock Guardrails: denied topics, PII redaction, filters
- Every prompt & response logged to an audit archive
- CloudWatch dashboards & Grafana observability
- Cost & resource controls built in

How a request flows through PAiTH



Section 2 — Continuum Design

An AI-driven modeling engine that turns ideas into working systems fast — capturing processes, generating architecture and code, and producing prototypes in hours. This is how each POV gets built.

DESIGN WITH CLARITY

Visual modeling

Event Storming, Domain-Driven Design, and Event Modeling create a shared understanding across business and engineering teams.

GENERATE WITH PRECISION

AI specs → code

Automatically produces requirements, data models, architecture, documentation, and working prototypes — ready for review and iteration.

DELIVER WITH CONFIDENCE

Export & deploy

Export to Jira or Azure DevOps, publish OpenAPI/CSV/JSON, and deploy into secure environments aligned with federal standards.

Hours

IDEA → PROTOTYPE

85%+

FASTER DEV TIME

You own

THE SOURCE CODE

CMMI L5

ISO · NIST 800-53

Available on the Tradewinds Solution Marketplace — pre-competed, readily awardable today. **Custom apps you fully own, eliminating expensive COTS licenses.**

Section 3 — SEC "Fed Claw" Fraud Monitor

The problem → the solution

Problem: Identifying complex pump-and-dump schemes across fragmented social media and trading data takes months — often after funds have moved offshore.

Solution: An autonomous engine ingests social sentiment (X, Reddit) and correlates it with real-time trading spikes to surface "Fed Claw" opportunities — freezing assets *before* they disappear.

Repository

 [secfedclaw-watch-v2](#) **CORE**

Quick-win prototype

A **60-day lookback** on a known 2025 manipulation case — proving the AI would have flagged the anomaly **~2 weeks earlier** than manual methods.

\$100M+

ANNUAL RECOVERY

-2 wks

TIME-TO-FLAG

"This isn't just surveillance — it's an asset-recovery engine that pays for itself by clawing back ill-gotten gains for defrauded investors."

Section 4 — Agentic Examiner

The problem → the solution

Problem: Examiners spend weeks gathering pre-exam intelligence on Registered Investment Advisers before arriving on-site.

Solution: An AI agent pulls a firm's ADV filings, past exam deficiencies, and news sentiment to build a "Risk Briefing Book" — and computes an explainable risk score the model narrates but does not invent.

Run it live

▶ SEC-POVs/agentic-exam-assistant · `./run.sh` → `localhost:8501`

Streamlit web app · Anthropic API for reasoning · live SEC EDGAR with offline fallback. Pick a firm → watch the agent work → render the briefing.

Quick-win prototype

A **45-day pilot** generating **10 automated Firm Risk Profiles** for upcoming RIA examinations.

24 hrs

PREP SAVED / EXAMINER

+15%

HIGH-RISK FINDINGS

"Our examiners should spend their time interviewing firm leadership, not digging through PDFs. This agent does the homework for them."

Protected by Zscaler + red-teamed with SPLX.ai

Zscaler enforces zero-trust access to the network and apps; SPLX.ai proves the AI itself holds up under attack. Defense in depth — from the connection to the model.

ZSCALER · ZERO-TRUST ACCESS

Secure the connection

- Zero-trust network access — users reach apps, never the network
- Full SSL inspection & cloud firewall on all traffic
- Inline DLP to stop sensitive filing & exam data exfiltration
- Identity-driven policy with continuous posture checks

SPLX.AI · AI RED-TEAMING

Secure the model

- Automated prompt-injection & jailbreak attacks
- Data-exfiltration & PII/PHI leakage probes
- Re-tests on every model or prompt change
- Scored, audit-ready findings mapped to AI risk frameworks

Layered defense: **Zscaler zero-trust** → **Cognito auth** → **Bedrock Guardrails** → **SPLX.ai red-team** → **audit archive**. Not a single checkpoint.

Section 6 — XBRL Quality Guard

The problem → the solution

Problem: Filers mis-tag XBRL data, making the SEC's economic models inaccurate.

Solution: A generative-AI auditor predicts the *expected* tag for each financial data point and flags human-tagged outliers for correction *before* they enter DERA models — a permanent 24/7 quality-control officer for the world's most important financial database.

Repositories

 continuum-xbrl-quality-guard **CORE**

 CD-sec-demo **DATABRICKS**

2-week proof of value

In **just 2 weeks**, scan the last **1,000 "Small Business" filings** and flag the top **5%** most-likely tagging errors for human review.

2 wks

TO FIRST RESULTS

90%

↑ DATA ACCURACY

Top 5%

ERRORS SURFACED

"Our policy is only as good as our data. This AI is a permanent, 24/7 quality-control officer for the world's most important financial database."

Business value across the program

One secure platform, built fast and proven safe — delivering measurable mission outcomes for three SEC divisions.

INITIATIVE	DIVISION	WHAT IT DELIVERS	BUSINESS VALUE
Fed Claw Fraud Monitor	Enforcement	Correlates social sentiment with trading spikes to freeze assets before they go dark	\$100M+ recovered / yr ~2 wks faster to flag
XBRL Quality Guard	DERA	Generative-AI auditor flags mis-tagged filings before they corrupt economic models	90% ↑ data accuracy top 5% errors surfaced
Agentic Exam Assistant	EXAMS	Builds a Risk Briefing Book from ADV filings, deficiencies, and news	24 hrs saved / examiner +15% high-risk findings
Continuum Design	All	AI modeling engine that builds each POV — specs, architecture, and prototypes in hours	85%+ faster dev customer-owned code
AWS PAiTH + Zscaler + SPLX.ai	All	Governed Bedrock platform, zero-trust access, and continuous AI red-teaming	Audit-ready trust defensible for ATO

Prove value in 45–60 days, then scale on a platform you already trust — three missions, one secure, governed, red-teamed AI foundation.